𝔐𝔈𝔯𝔠𝔲𝔯𝔶: A multilinear Polynomial Commitment Scheme with constant proof size and no prover FFTs

Liam Eagen

Alpen Labs

Ariel Gabizon

Aztec Labs

February 27, 2025

Abstract

We construct a pairing-based polynomial commitment scheme for multilinear polynomials of size n where constructing an opening proof requires $O(n)$ field operations, and $2n + O(\sqrt{n})$ scalar multiplications. Moreover, the opening proof consists of a constant number of field elements. This is a significant improvement over previous works which would require either

1. $O(n \log n)$ field operations; or

2. $O(\log n)$ size opening proof.

The main technical component is a new method of verifiably folding a witness via univariate polynomial division. As opposed to previous methods, the proof size and prover time remain constant *regardless of the folding factor*.

1 Introduction

Polynomial Commitment Schemes (PCSs)[KZG10] allow a party to commit to a polynomial and later prove an evaluation of the polynomial is correct. That is, for a commitment cm and values a, b; a prover **P** can produce a proof that $\mathsf{cm} = \mathsf{com}(f(X))$ and $f(a) = b$. PCSs form an essential part of most modern Succinct Non-interactive ARguments of Knowledge (SNARKs). They allow a protocol designer to focus on designing a so-called polynomial Interactive Oracle Proof which can then be compiled, via a PCS, to a SNARK (see [BFS19, GWC19, CHM+19] for descriptions of such compilers). In fact,

1

many of the most important properties of a SNARK, like proof size, verifier complexity, and cryptographic assumptions, follow primarily from the PCS. The earliest polynomial commitment schemes [KZG10] supported univariate polynomials and were used to construct SNARKs like Plonk [GWC19] and Marlin [CHM+19] with $O(n \log n)$ prover complexity and constant proof size. A different class of SNARKs [Set19, CBBZ22] arising from the sumcheck protocol [LFKN92] have linear prover time, but require *multilinear* Polynomial Commitment Schemes (ml-PCS's).

1.1 Our results

Existing transformations from a univariate PCS to ml-PCS are either linear time but require a logarithmic opening proof size, like gemini [BCHO22] and zeromorph [KT23], or have constant size opening proofs but incur an additional $O(n \log n)$ prover cost to perform univariate polynomial multiplication via FFT's. We propose a new protocol that goes beyond this tradeoff: a protocol with constant proof size and only $O(n)$ prover operations (besides the $O(\lambda n/(\log(\lambda n))$ operations for muti-scalar multiplications arising in KZG commitments). Additionally, this is concretely more efficient than existing schemes, with a smaller constant factor of prover work ($1 + o(1)$ or $2 + o(1)$ compared to 2.5 for zeromorph) and concretely small proof sizes.

Table 1: Comparison of pairing-based ml-PCS. $\mathbb{G}$ denotes a scalar multiplication. All verifiers below additionally require two pairings.

Scheme	Proof size	Prover Work	Verifier Work
univariate-based e.g.[PH23]	$O(1)\ \mathbb{F}$	$O(n \log n)\ \mathbb{F}, O(n)\ \mathbb{G}$	$O(\log n)\ \mathbb{F}, O(1)\ \mathbb{G}$
gemini [BCHO22]	$O(\log n)\ \mathbb{F}$	$O(n)\ \mathbb{F}, 3n\ \mathbb{G}$	$O(\log n)\ \mathbb{F}, O(\log n)\ \mathbb{G}$
zeromorph [KT23]	$O(\log n)\ \mathbb{F}$	$O(n)\ \mathbb{F}, 2.5n\ \mathbb{G}$	$O(\log n)\ \mathbb{F}, O(\log n)\ \mathbb{G}$
𝔐𝔢𝔯𝔠𝔲𝔯𝔶 (this work)	$O(1)\ \mathbb{F}$	$O(n)\ \mathbb{F}, 2n + O(\sqrt{n})\ \mathbb{G}$	$O(\log n)\ \mathbb{F}, O(1)\ \mathbb{G}$

2 Overview of technique

Remark 2.1. *In this overview, we use some of the notation defined in Sections 3.1 and 3.2.*

Our technique is best thought of as an improvement of the gemini ml-PCS [BCHO22]. Let's start by recalling how gemini works. gemini commits to a multilinear function as a *univariate* KZG commitment[KZG10]. Specifically, fix a vector $f \in \mathbb{F}^n$ describing the function's values on the boolean cube B_s where $s = \log n$. That is, we think of f as representing the multilinear

$$M(X_0, \ldots, X_{s-1}) = \sum_{i<n} \mathbf{eq}(i, X_0, \ldots, X_{s-1}) f_i.$$

2

(Here, as explained in Section 3.2, we interpret i as its binary decomposition $(i_0, \ldots, i_{s-1})$ when used as input to **eq**.) Let $\mathsf{srs} = \left\{ \left[x^i \right]_1 \right\}_{i<n}$ be a KZG structured reference string. gemini outputs $\mathsf{cm} = [f(x)]_1 = \sum_{i<n} f_i \left[x^i \right]_1$ as a commitment to M.

Now suppose prover $\mathbf{P}$ wants to convince verifier $\mathbf{V}$ that $M(u) = v$, for some $u = (u_0, \ldots, u_{s-1}) \in \mathbb{F}^s$. In gemini, $\mathbf{P}$ sends commitments $\mathsf{cm}_1, \ldots, \mathsf{cm}_s$ to the s incremental restrictions leading to evaluation at u. Namely, to $M_1 = M(u_0, X_1, \ldots, X_{s-1})$, $M_2 = M(u_0, u_1, X_2, \ldots, X_{s-1}), \ldots, M_s = M(u_0, \ldots, u_{s-1})$. Assuming $\mathbf{P}$ sent commitments to the correct functions, all that is needed is to check that cm_s is the commitment to the constant v. Of course, the interesting part is proving the commitments *are* to the correct functions!

For this purpose, gemini exploits a connection between M and its corresponding univariate $f(X)$: Write $f(X) = f_0(X^2) + X f_1(X^2)$, for $f_0(X), f_1(X)$ of degree $< n/2$. Let $f_{u_1}(X)$ be the univariate corresponding to M_1 defined above. Then, we have

$$f_{u_1}(X) = (1 - u_1) f_0(X) + u_1 f_1(X).$$

Additionally, we can evaluate f_0 and f_1 via f using the equations

$$f_0(X^2) = \frac{f(X) + f(-X)}{2}, f_1(X^2) = \frac{f(X) - f(-X)}{2X}$$

Thus, we can perform consistency checks between each pair $\mathsf{cm}_{i-1}, \mathsf{cm}_i$, via univariate KZG openings at a random challenge, inductively showing cm_i is indeed the commitment to the next desired restriction. Of course, we get $O(s) = O(\log n)$ proof length due to this sequence of restriction commitments.

Here is a first idea on how to reduce proof length. Protocols based on univariate polynomials allow us to do multilinear evaluation in $O(n \log n)$ prover time with constant proof size (e.g. Section 5 of [PH23]) . Choose a parameter t and set $b = 2^t$. We can run *only the first t rounds* of gemini, reaching a restricted multilinear on $n - t$ variables. If $n' = n/b \leq n/\log n$, we can afford to run a univariate protocol with $O(n' \log n') = O(n)$ prover time to evaluate $M_t(u_t, \ldots, u_{s-1})$. This still doesn't take us to overall constant proof size - as we need to use a super-constant t to reach such n'. (For us $t = \log n / 2$ will be optimal, although $t \geq \log \log n$ suffices here.)

This raises the question - can we "skip" the intermediate gemini rounds and send *only* the commitment cm_t, and directly prove it is consistent with the original cm? Extrapolating the gemini strategy in the natural way, we get the answer - yes, but not with constant proof size: We can decompose f into b polynomials of degree $< n/b$: $f(X) = \sum_{0 \leq i < b} X^i f_i(X^b)$. As in the $b = 2$ case, one can show the univariate $f'(X)$ corresponding to M_t is a linear combination of the $\{f_i(X)\}$. Moreover, evaluating the f_i using f (for the consistency check) can be done. However, it requires b evaluations of f. Specifically, $f_i(r^b)$ is a linear combination of $\left\{ f(r), f(r\omega), \ldots, f(r\omega^{b-1}) \right\}$ where ω is a primitive b'th root of unity.

Our central innovation is a different way to prove cm_t is correct *with* constant proof size. The (univariate corresponding to the) correct restricted polynomial is

$$h(X) = \sum_{0 \leq i < b} \mathbf{eq}(i, u_1) f_i(X).$$

3

Let $g(X) := f(X) \mod X^b - \alpha$. Calculation shows

$$g(X) = \sum_{0 \leq i < b} X^i f_i(\alpha).$$

The multilinear $\hat{g}$ corresponding to $g(X)$ is

$$\hat{g}(X_0, \ldots X_{t-1}) = \sum_{0 \leq i < b} \mathbf{eq}(i, X_0, \ldots, X_{t-1}) f_i(\alpha)$$

In particular, we have

$$\hat{g}(u_1) = \sum_{0 \leq i < b} \mathbf{eq}(i, u_1) f_i(\alpha) = h(\alpha).$$

In words, the evaluation of g at u_1 as a *multilinear* corresponds to the evaluation of h at α as a univariate! We can use standard univariate KZG to open cm_t at α. And, crucially, we can afford to evaluate $\hat{g}(u_1)$ using the aforementioned univariate protocols as it is of size b rather than n. In summary, we can show a committed polynomial corresponds to the correct restriction. And now, again, we can afford to open h as a multilinear at u_2 using univariate protocols as it has size n/b rather than n.

Comparison to sumcheck   It is instructive to see what happens if we try to get a similar result via a modification of the sumcheck protocol [LFKN92]. Note first that a multilinear evaluation can indeed written as a sum over the function's values on B_s multiplied by the **eq** function:

$$M(u) = \sum_{b \in B_s} \mathbf{eq}(b, u) M(b).$$

The classic sumcheck protocol, like gemini, works by $\log n$ reductions of the domain size by a factor of two; each round fixing one more variable of the summed function. In the above spirit, we could look at a modified sumcheck protocol, where the first variable ranges over a domain of super-constant size b. The first round univariate P_1 would thus have degree roughly $2b$. To maintain constant proof length, we could send a commitment to P_1 rather than its coefficients (as usually done in sumcheck). However, *computing P_1* would require superlinear time $O(n \log b)$ - as we need to perform a b-size FFT for n/b values of the second variable appearing in the sum.

3   Preliminaries

3.1   Terminology and conventions

We work with integer parameter n that we'll assume throughout the paper is of the form $n = 2^{2t}$ for integer $t > 0$. We'll denote its square root by $b := 2^t = \sqrt{n}$. We index vectors starting from zero. For example, for $g \in \mathbb{F}^b$ we have $g = (g_0, \ldots, g_{b-1})$.

We associate vectors with univariate polynomials in the following natural way: Given $g \in \mathbb{F}^b$ we denote $g(X) := \sum_{0 \le i < b} g_i X^i$.

We make the convention that integer ranges in sums begin at zero if not specified otherwise. Thus, we write $g(X) = \sum_{i<b} g_i X^i$.

We assume vectors of size n are indexed by two indices ranging over $\{0, \ldots, b-1\}$. Thus, for $f \in \mathbb{F}^n$, we have $f = (f_{0,0}, \ldots, f_{0,b-1}, \ldots, f_{b-1,0}, \ldots, f_{b-1,b-1})$. Accordingly, for $0 \le i < b$, we denote by f_i the vector $(f_{i,0}, \ldots, f_{i,b-1})$.

In particular, for $f \in \mathbb{F}^n$ we have under these notations that

$$f(X) := \sum_{i<b} X^i f_i(X^b) = \sum_{i<b} \sum_{j<b} f_{i,j} X^{i+j \cdot b}$$

We denote by B_t the binary cube $\{0,1\}^t$ of dimension t.

3.2   Multilinear polynomials

Let $n = 2^{2t}$. We define the well-known **eq** multilinear polynomial in $4t$ variables.

$$\mathbf{eq}(x, y) := \prod_{i=0}^{t-1} (x_i y_i + (1 - x_i)(1 - y_i))$$

We have for $x, y \in B_{2t}$, $\mathbf{eq}(x,y) = 1$ when $x = y$ and $\mathbf{eq}(x,y) = 0$ otherwise.

We use the convention that an integer $0 \le i < n$ can be used as an input to **eq** by interpreting i as its binary representation. Namely, for $0 \le i < n$, $u \in \mathbb{F}^t$, $\mathbf{eq}(i, u) := \mathbf{eq}(i_1, \ldots, i_t, u)$ where $i = \sum_{j \in [t]} i_j 2^{j-1}$.

For $a \in \mathbb{F}^n$, we define $\hat{f}$ to be the multilinear polynomial obtaining f's values on the boolean cube. Namely,

$$\hat{a}(X_0, \ldots, X_{t-1}) := \sum_{i<n} \mathbf{eq}(i, X_1, \ldots, X_t) \cdot a_i.$$

3.3   The algebraic group model

We introduce some terminology from [GWC19] to capture analysis in the Algebraic Group Model of Fuchsbauer, Kiltz and Loss[FKL18].

In our protocols, by an *algebraic adversary* $\mathcal{A}$ in an SRS-based protocol we mean a poly(λ)-time algorithm which satisfies the following.

- For $i \in \{1, 2\}$, whenever $\mathcal{A}$ outputs an element $A \in \mathbb{G}_i$, it also outputs a vector v over $\mathbb{F}$ such that $A = < v, \mathsf{srs_i} >$.

First we say our srs *has degree* Q if all elements of $\mathsf{srs_i}$ are of the form $[f(x)]_i$ for $f \in \mathbb{F}_{<Q+1}[X]$ and uniform $x \in \mathbb{F}$. In the following discussion let us assume we are executing a protocol with a degree Q SRS, and denote by $f_{i,j}$ the corresponding polynomial for the j'th element of $\mathsf{srs_i}$.

Denote by a, b the vectors of $\mathbb{F}$-elements whose encodings in $\mathbb{G}1, \mathbb{G}2$ an algebraic adversary $\mathcal{A}$ outputs during a protocol execution; e.g., the j'th $\mathbb{G}1$ element output by $\mathcal{A}$ is $[a_j]_1$.

By a "real pairing check" we mean a check of the form

$$(a \cdot T_1) \cdot (T_2 \cdot b) = 0$$

for some matrices T_1, T_2 over $\mathbb{F}$. Note that such a check can indeed be done efficiently given the encoded elements and the pairing function $e : \mathbb{G}1 \times \mathbb{G}2 \to \mathbb{G}_t$.

Given such a "real pairing check", and the adversary $\mathcal{A}$ and protocol execution during which the elements were output, define the corresponding "ideal check" as follows. Since $\mathcal{A}$ is algebraic when he outputs $[a_j]_i$ he also outputs a vector v such that, from linearity, $a_j = \sum v_\ell f_{i,\ell}(x) = R_{i,j}(x)$ for $R_{i,j}(X) := \sum v_\ell f_{i,\ell}(X)$. Denote, for $i \in \{1, 2\}$ the vector of polynomials $R_i = (R_{i,j})_j$. The corresponding ideal check, checks as a polynomial identity whether

$$(R_1 \cdot T_1) \cdot (T_2 \cdot R_2) \equiv 0$$

The following lemma is inspired by [FKL18]'s analysis of [Gro16], and tells us that for soundness analysis against algebraic adversaries it suffices to look at ideal checks. Before stating the lemma we define the Q-DLOG assumption similarly to [FKL18].

Definition 3.1. *Fix integer Q. The Q-DLOG assumption for $(\mathbb{G}1, \mathbb{G}2)$ states that given*

$$[1]_1, [x]_1, \ldots, [x^Q]_1, [1]_2, [x]_2, \ldots, [x^Q]_2$$

for uniformly chosen $x \in \mathbb{F}$, the probability of an efficient $\mathcal{A}$ outputting x is $\mathsf{negl}(\lambda)$.

Lemma 3.2. *Assume the Q-DLOG for $(\mathbb{G}1, \mathbb{G}2)$. Given an algebraic adversary $\mathcal{A}$ participating in a protocol with a degree Q SRS, the probability of any real pairing check passing is larger by at most an additive $\mathsf{negl}(\lambda)$ factor than the probability the corresponding ideal check holds.*

See [GWC19] for the proof.

3.4 Polynomial commitment schemes for multilinear polynomials

We give a formal definition of a ml-PCS secure in the algebraic group model.

Definition 3.3. *Let $n = 2^t$. A multilinear polynomial commitment scheme (ml-PCS) consists of*

- $\mathsf{gen}(n)$ - *a randomized algorithm that outputs an SRS* srs *0.*

- $\mathsf{com}(f, \mathsf{srs})$ - *that given a polynomial $f \in \mathbb{F}^n$ returns a commitment cm to f.*

- *A public coin protocol $\mathsf{open}(\mathsf{cm}, n, u, v)$ between parties $\mathbf{P}$ and $\mathbf{V}$. $\mathbf{P}$ is given $f \in \mathbb{F}^n$. $\mathbf{P}$ and $\mathbf{V}$ are both given integer n, cm- the purported commitment to f, $u \in \mathbb{F}^t$ and $v \in \mathbb{F}$ - the purported value $\hat{f}(u)$.*

6

such that

- **Completeness:** *Suppose that ,* $\mathsf{cm} = \mathsf{com}(f, \mathsf{srs})$*. Then if* open *is run correctly with values* $n, \mathsf{cm}, u, v = \hat{f}(u)$*,* **V** *outputs* $accept$ *with probability one.*

- **Knowledge soundness in the algebraic group model:** *There exists an efficient E such that for any algebraic adversary* $\mathcal{A}$ *the probability of* $\mathcal{A}$ *winning the following game is* $\mathsf{negl}(\lambda)$ *over the randomness of* $\mathcal{A}$ *and* gen*.*

  1. *Given* srs*,* $\mathcal{A}$ *outputs* n, cm*.*
  2. *E, given access to the messages of* $\mathcal{A}$ *during the previous step, outputs* $f \in \mathbb{F}^n$*.*
  3. $\mathcal{A}$ *outputs* $u \in \mathbb{F}^t$*,* $v \in \mathbb{F}$*.*
  4. $\mathcal{A}$ *takes the part of* **P** *in the protocol* open *with inputs* n, cm, u, v*.*
  5. $\mathcal{A}$ *wins if*
     - **V** *outputs* $accept$ *at the end of the protocol.*
     - $\hat{f}(u) \neq v$*.*

4 Components

In this section we go over known components (with some new optimizations), that will be used in our main protocol in Section 6. The treatment will be semi-formal, and assume basic familiarity with the KZG polynomial commitment scheme [KZG10]. The formal treatment will be part of the description and knowledge soundness proof of the main protocol in Section 6.

4.1 Inner products in $O(b \log b)$ time.

Fix polynomials $g_1(X) = \sum_{i=0}^{d_1} a_i X^i, g_2 = \sum_{i=0}^{d_2} b_i X^i$ in $\mathbb{F}[X]$. We define $< g_1, g_2 >$ to be $\sum_{i=0}^{d} a_i b_i$ where $d := \min \{d_1, d_2\}$. We present a convenient way to verify inner products $< g_1, g_2 >$ similar to [BCC$^+$16, MBKM19]. The basic observation is that $< g_1, g_2 >$ is the constant coefficient of the rational function $R(X) := g_1(X)g_2(1/X)$. Thus, $< g_1, g_2 >= v$ is equivalent to the existence of polynomials S_1, S_2 such that

$$g_1(X)g_2(1/X) = 1/X \cdot S_1(1/X) + v + X \cdot S_2(X).$$

We can thus sends commitments to S_1, S_2 as proof of the correctness of v. As an optimization, we observe that we can "symmetrize" R an look instead at the rational function

$$R'(X) := g_1(X)g_2(1/X) + g_1(1/X)g_2(1/X).$$

The advantage of R' is that the negative and positive coefficients are equal. Thus, $< g_1, g_2 >= v$ is equivalent to the existence of $S(X) \in \mathbb{F}[X]$ such that

$$g_1(X)g_2(1/X) + g_1(1/X)g_2(1/X) = 2v + X \cdot S(X) + (1/X)S(1/X).$$

Claim 4.1. *Suppose $g_1(X), g_2(X) \in \mathbb{F}_{<b}[X]$. Let $S(X)$ be as defined above. Then S can be computed in $O(b \log b)$ $\mathbb{F}$-operations.*

Proof. When $g_1(X), g_2(X) \in \mathbb{F}_{<b}[X]$ we multiply the equation above by X^{b-1} to get

$$X^{b-1}(g_1(X)g_2(1/X) + g_1(1/X)g_2(X)) = X^{b-1}(2v + X \cdot S(X) + (1/X)S(1/X)).$$

We can use an $O(b \log b)$ time FFT to evaluate the LHS on $2b$ points. We then do an inverse FFT to get the coefficients $c_0, \ldots, c_{2b-2}$ of the LHS. Now, we can output $S = (c_b, \ldots, c_{2b-2})$. $\qquad \square$

Batching inner product checks   Suppose we now have two inner product claims $< g_1, g_2 >= v_1$ and $< h_1, h_2 >= v_2$. If we choose random $\gamma \in \mathbb{F}$ with high probability

Claim 4.2. *Fix polynomials $g_1, g_2, h_1, h_2 \in \mathbb{F}[X]$ and $v_1, v_2 \in \mathbb{F}$. Suppose $< g_1, g_2 > \neq v_1$ or $< h_1, h_2 > \neq v_2$. Then, e.w.p $1/|\mathbb{F}|$ over $\gamma \in \mathbb{F}$ there does not exist $S(X) \in \mathbb{F}[X]$ such that*

$$g_1(X)g_2(1/X) + g_1(1/X)g_2(X) + \gamma(h_1(X)h_2(1/X) + h_1(1/X)h_2(X))$$

$$= 2(v_1 + \gamma v_2) + X \cdot S(X) + (1/X)S(1/X)$$

Proof. Denote $(v_1', v_2') = (< g_1, g_2 >, < h_1, h_2 >))$. The constant coefficient of the LHS is $v_1' + \gamma v_2'$ To satisfy the equation in the claim for some S, we need

$$v_1' + \gamma v_2' = v_1 + \gamma v_2,$$

which can hold for at most one γ when $(v_1, v_2) \neq (v_1', v_2')$. $\qquad \square$

4.2 Multilinear evaluations as inner products of univariate polynomials.

For $u \in B_t$ define the polynomial $P_u(X) := \sum_{i<b} \mathbf{eq}(i, u)X^i$. Note that for $g(X) \in \mathbb{F}_{<b}[X]$, we have

$$< P_u, g >= \sum_{i<b} \mathbf{eq}(i, u)g_i = \hat{g}(u).$$

As leveraged in [BGH19], we have the product formula

$$P_u(X) = \prod_{i=0}^{t-1} \left(u_i X^{2^i} + 1 - u_i \right),$$

implying $P_u(X)$ can be evaluated in $O(t)$ $\mathbb{F}$-operations.

Using the polynomials P_u, mutlilinear evaluations can be proven in a batched manner based on Claim 4.2: Suppose we want to show given committed univariates $g, h \in \mathbb{F}[X]$, $u_1, u_2 \in \mathbb{F}^t, v_1, v_2 \in \mathbb{F}$ that $\hat{g}(u_1) = v_1$ and $\hat{h}(u_2) = v_2$.

1. **V** sends random $\gamma \in \mathbb{F}$.

2. **P** sends commitment to S such that

$$g(X)P_{u_1}(1/X) + g(X)P_{u_1}(1/X) + \gamma(h(X)P_{u_2}(1/X) + h(1/X)P_{u_2}(X))$$

$$= 2(v_1 + \gamma v_2) + X \cdot S(X) + (1/X)S(1/X).$$

3. **V** chooses a random $\mathfrak{z} \in \mathbb{F}$.

4. **P** sends and proves correctness of the values of g, h and S on $\mathfrak{z}, 1/\mathfrak{z}$.

5. **V** evaluates P_{u_1}, P_{u_2} at $\mathfrak{z}, 1/\mathfrak{z}$.

6. **V** checks the above equation holds at $\mathfrak{z}$.

4.3   Degree checks

The idea presented here is from [Tha23]. Suppose **P** wants to prove to **V** that cm is a commitment to a polynomial $g \in \mathbb{F}_{<b}[X]$. Let $D(X) := X^{b-1}g(1/X)$. The idea is that D is a polynomial if and only if $g(X)$ has degree $< b$. Thus, assuming our structured reference string doesn't contain negative powers, **P** can commit to D if and only if $g(X) \in \mathbb{F}_{<b}[X]$.

This motivates the following protocol.

1. **P** sends a commitment d to $D(X)$.

2. **V** chooses random $\mathfrak{z} \in \mathbb{F}$.

3. **P** sends $D_{\mathfrak{z}} := D(\mathfrak{z}), \bar{g}_{\mathfrak{z}} := g(1/\mathfrak{z})$, and uses KZG to prove their correctness.

4. **V** can now check D's correctness on $\mathfrak{z}$, using the equation

$$D_{\mathfrak{z}} \stackrel{?}{=} \mathfrak{z}^{b-1}\bar{g}_{\mathfrak{z}}.$$

5   Univariate division

Our protocol crucially relies on the following simple claim about division by a polynomial of the form $X^b - \alpha$.

Claim 5.1. *Fix integers $b > 0$ and let $n = b^2$. Fix $\alpha \in \mathbb{F}$, and $f(X) \in \mathbb{F}_{<n}[X]$. Let $f_0(X), \ldots, f_{b-1}(X) \in \mathbb{F}_{<b}[X]$ be such that $f(X) = \sum_{i<b} X^i f_i(X^b)$. Let $g(X) \in \mathbb{F}_{<b}[X], q(X) \in \mathbb{F}[X]$ be such that*

$$f(X) = (X^b - \alpha) \cdot q(X) + g(X).$$

Then,

1. $g(X) = \sum_{i<b} X^i f_i(\alpha)$.

9

2. *The coefficients of $q(X)$ can be computed in $O(n)$ $\mathbb{F}$-operations.*

Proof. To see the first item, note that reduction mod $X^b - \alpha$ corresponds to substituting α into X^b inside each $f_i(X^b)$ in the expression $\sum_{i<b} X^i f_i(X^b)$. We proceed to the computation of $q(X)$. We compute for each $0 \le i < b$, the coefficients of the quotient $q_i(X) \in \mathbb{F}[X]$ such that
$$f_i(X) = q_i(X)(X - \alpha) + f_i(\alpha).$$

Using Horner's method for division by the linear polynomial $X - \alpha$ this requires only n multiplications and additions in $\mathbb{F}$. Now, we have that

$$f(X) = \sum_{i<b} X^i f_i(X^b) = \sum_{i<b} X^i \left(q_i(X^b)(X^b - \alpha) + f_i(\alpha) \right) = q(X)(X^b - \alpha) + g(X),$$

for $q(X) := \sum_{i<b} X^i q_i(X^b)$. Thus, the coefficients of $q(X)$ are simply the interleaving of the coefficients of the $\{q_i(X)\}$. $\qquad\square$

6   Main Construction

$\mathfrak{MERCURY}$ is the tuple $(\mathsf{gen}, \mathsf{com}, \mathsf{open})$ described next.

$\underline{\mathsf{gen}(n)}$: Choose random $x \in \mathbb{F}$ and outputs $\left\{ [1]_1, [x]_1, \ldots, [x^{n-1}]_1, [1]_2, [x]_2 \right\}$

$\underline{\mathsf{com}(n, f, \mathsf{srs})}$: Output $\sum_{i<b} \sum_{j<b} f_{i,j} \cdot [x^{i \cdot b + j}]_1$.

$\underline{\mathsf{open}(n, \mathsf{cm}, u, v; f)}$:

1. Committing to partial sums:

    (a) $\mathbf{P}$ computes the polynomial to $h(X) := \sum_{i<b} \mathbf{eq}(i, u_1) f_i(X)$. Note that the coefficient of X^j in $h(X)$ is $\sum_{i<b} \mathbf{eq}(i, u_1) f_{i,j}$ - hence we think of it as a commitment to partial sums.

    (b) $\mathbf{P}$ computes and sends $\mathsf{h} := [h(x)]_1$.

2. Committing to "folded" polynomial g:

    (a) $\mathbf{V}$ sends random $\alpha \in \mathbb{F}$.

    (b) $\mathbf{P}$ computes polynomials $g(X) \in \mathbb{F}_{<b}[X]$ and $q(X) \in \mathbb{F}[X]$ such that

    $$f(X) = (X^b - \alpha) \cdot q(X) + g(X).$$

    (c) $\mathbf{P}$ computes and sends $\mathsf{q} := [q(x)]_1$ and $\mathsf{g} := [g(x)]_1$.

10

3. Sending proofs of correctness for h and the degree of g:

   (a) **V** sends a random batching challenge $\gamma \in \mathbb{F}$.

   (b) **P** computes and sends $\mathsf{s} = [S(x)]_1$ where $S(X) \in \mathbb{F}[X]$ is such that

   $$g(X)P_{u_1}(1/X) + g(1/X)P_{u_1}(X) + \gamma \cdot (h(X)P_{u_2}(1/X) + h(1/X)P_{u_2}(X))$$

   $$= h(\alpha) + \gamma \cdot v + X \cdot S(X) + (1/X)S(1/X).$$

   (c) **P** computes and sends $\mathsf{d} := [D(x)]_1$ where

   $$D(X) := X^{b-1}g(1/X).$$

4. KZG evaluations:

   (a) **V** sends a random evaluation challenge $\mathfrak{z} \in \mathbb{F}$.

   (b) **P** sends the values $f_\mathfrak{z} := f(\mathfrak{z}), q_\mathfrak{z} := q(\mathfrak{z}), g_\mathfrak{z} := g(\mathfrak{z}), \bar{g}_\mathfrak{z} := g(1/\mathfrak{z}), h_\mathfrak{z} := h(\mathfrak{z}), \bar{h}_\mathfrak{z} := h(1/\mathfrak{z}), h_\alpha := h(\alpha), s_\mathfrak{z} := S(\mathfrak{z}), \bar{s}_\mathfrak{z} := S(1/\mathfrak{z})$.

   (c) **V** computes the expected value for $D(\mathfrak{z})$ - $D_\mathfrak{z} := \mathfrak{z}^{b-1}\bar{g}_\mathfrak{z}$.

   (d) **V** sends a random KZG batching challenge $\eta \in \mathbb{F}$.

   (e) **P** computes and sends the KZG opening proof $\pi_\mathfrak{z}$ for the values $f_\mathfrak{z}$ and $q_\mathfrak{z}$. That is $\pi_\mathfrak{z} := [H(x)]_1$ for

   $$H(X) := \frac{f(X) - f(\mathfrak{z}) + \eta(q(X) - q(\mathfrak{z}))}{X - \mathfrak{z}}.$$

   (f) **P** computes and send a batched KZG opening proof π' for the rest of the values sent in step 4b, as described in Section 4 of [BDFG20].

   (g) **V** checks the proof $\pi_\mathfrak{z}$ as in [KZG10]:

   $$e(\mathsf{cm} - [f_\mathfrak{z}]_1 + \eta(\mathsf{q} - [q_\mathfrak{z}]_1), [1]_2) = e(\pi_\mathfrak{z}, [x]_2).$$

   (h) **V** checks the opening proof π' as described in [BDFG20].

   (i) **V** checks the equation

   $$g_\mathfrak{z}P_{u_1}(1/\mathfrak{z}) + \bar{g}_\mathfrak{z}P_{u_1}(\mathfrak{z}) + \gamma(h_\mathfrak{z}P_{u_2}(1/\mathfrak{z}) + \bar{h}_\mathfrak{z}P_{u_2}(\mathfrak{z})) = h_\alpha + \gamma v + \mathfrak{z}s_\mathfrak{z} + (1/\mathfrak{z})\bar{s}_\mathfrak{z}.$$

   (j) If one of the checks in steps 4g-4i fails, **V** outputs *reject*. Otherwise **V** outputs *accept*.

Runtime of **P**: Computing $q(X)$ in step 2b requires $O(n)$ operations by Claim 5.1. Computing q and $\pi_{\mathfrak{z}}$ requires two MSMs of size n. All other steps are on polynomials of size $O(b) = O(\sqrt{n})$. Thus other commitments clearly require $O(\sqrt{n})$ scalar multiplications. It is easy to see other steps require $o(n)$ $\mathbb{F}$-operations. The least trivial of these is perhaps the computation of $S(X)$ shown to require $O(b \log b) = o(n)$ operations in Claim 4.1.

Proving knowledge soundness: Let $\mathcal{A}$ be an efficient algebraic adversary participating in the Knowledge Soundness game from Definition 3.3. We show its probability of winning the game is $\mathsf{negl}(\lambda)$. Let $f \in \mathbb{F}^n$ be the vector sent by $\mathcal{A}$ in the third step of the game such that $\mathsf{cm} = [f(x)]_1$. As $\mathcal{A}$ is algebraic, when sending the commitments h,q,g,s,d,$\pi_{\mathfrak{z}}$,π' during protocol execution it also sends polynomials $h(X), q(X), g(X), S(X), D(X), H(X), Q(X) \in \mathbb{F}_{<n}[X]$ such that the former are their corresponding commitments. Let E be the event that **V** outputs accept. Note that the event that $\mathcal{A}$ wins the knowledge soundness game is contained in E. E implies all pairing checks have passed. Let $A \subset E$ be the event that one of the corresponding ideal pairing checks as defined in Section 3.3 didn't pass. According to Lemma 3.2, $\mathsf{prob}(A) = \mathsf{negl}(\lambda)$.

References

[BCC+16]   J. Bootle, A. Cerulli, P. Chaidos, J. Groth, and C. Petit. Efficient zero-knowledge arguments for arithmetic circuits in the discrete log setting. pages 327–357, 2016.

[BCHO22]   J. Bootle, A. Chiesa, Y. Hu, and M. Orrù. Gemini: Elastic snarks for diverse environments. *IACR Cryptol. ePrint Arch.*, page 420, 2022.

[BDFG20]   D. Boneh, J. Drake, B. Fisch, and A. Gabizon. Efficient polynomial commitment schemes for multiple points and polynomials. *IACR Cryptol. ePrint Arch.*, page 81, 2020.

[BFS19]   B. Bünz, B. Fisch, and A. Szepieniec. Transparent snarks from DARK compilers. *IACR Cryptol. ePrint Arch.*, page 1229, 2019.

[BGH19]   S. Bowe, J. Grigg, and D. Hopwood. Halo: Recursive proof composition without a trusted setup. *IACR Cryptol. ePrint Arch.*, page 1021, 2019.

[CBBZ22]   B. Chen, B. Bünz, D. Boneh, and Z. Zhang. Hyperplonk: Plonk with linear-time prover and high-degree custom gates. *IACR Cryptol. ePrint Arch.*, page 1355, 2022.

[CHM+19]   A. Chiesa, Y. Hu, M. Maller, P. Mishra, N. Vesely, and N. P. Ward. Marlin: Preprocessing zksnarks with universal and updatable SRS. *IACR Cryptology ePrint Archive*, 2019:1047, 2019.

[FKL18]   G. Fuchsbauer, E. Kiltz, and J. Loss. The algebraic group model and its applications. In *Advances in Cryptology - CRYPTO 2018 - 38th Annual International Cryptology Conference, Santa Barbara, CA, USA, August 19-23, 2018, Proceedings, Part II*, pages 33–62, 2018.

[Gro16]   J. Groth. On the size of pairing-based non-interactive arguments. In *Advances in Cryptology - EUROCRYPT 2016 - 35th Annual International Conference on the Theory and Applications of Cryptographic Techniques, Vienna, Austria, May 8-12, 2016, Proceedings, Part II*, pages 305–326, 2016.

[GWC19]   A. Gabizon, Z. J. Williamson, and O. Ciobotaru. PLONK: permutations over lagrange-bases for oecumenical noninteractive arguments of knowledge. *IACR Cryptology ePrint Archive*, 2019:953, 2019.

[KT23]   T. Kohrita and P. Towa. Zeromorph: Zero-knowledge multilinear-evaluation proofs from homomorphic univariate commitments. *IACR Cryptol. ePrint Arch.*, page 917, 2023.

[KZG10]   A. Kate, G. M. Zaverucha, and I. Goldberg. Constant-size commitments to polynomials and their applications. pages 177–194, 2010.

[LFKN92]   C. Lund, L. Fortnow, H. J. Karloff, and N. Nisan. Algebraic methods for interactive proof systems. *J. ACM*, 39(4):859–868, 1992.

[MBKM19]   M. Maller, S. Bowe, M. Kohlweiss, and S. Meiklejohn. Sonic: Zero-knowledge snarks from linear-size universal and updateable structured reference strings. *IACR Cryptology ePrint Archive*, 2019:99, 2019.

[PH23]   S. Papini and U. Haböck. Improving logarithmic derivative lookups using GKR. *IACR Cryptol. ePrint Arch.*, page 1284, 2023.

[Set19]   S. T. V. Setty. Spartan: Efficient and general-purpose zksnarks without trusted setup. *IACR Cryptol. ePrint Arch.*, page 550, 2019.

[Tha23]   S. Thakur. A flexible snark via the monomial basis. *IACR Cryptol. ePrint Arch.*, page 788, 2023.