𝔐𝔈𝔕ℭ𝔘ℜ𝔜: A multilinear Polynomial Commitment Scheme with constant proof size and no prover FFTs

Liam Eagen

Alpen Labs

Ariel Gabizon

Aztec Labs

March 5, 2025

Abstract

We construct a pairing-based polynomial commitment scheme for multilinear polynomials of size n where constructing an opening proof requires $O(n)$ field operations, and $2n + O(\sqrt{n})$ scalar multiplications. Moreover, the opening proof consists of a constant number of field elements. This is a significant improvement over previous works which would require either

1. $O(n \log n)$ field operations; or

2. $O(\log n)$ size opening proof.

The main technical component is a new method of verifiably folding a witness via univariate polynomial division. As opposed to previous methods, the proof size and prover time remain constant *regardless of the folding factor*.

1 Introduction

Polynomial Commitment Schemes (PCSs)[KZG10] allow a party to commit to a polynomial and later prove an evaluation of the polynomial is correct. That is, for a commitment cm and values a, b; a prover **P** can produce a proof that $\mathsf{cm} = \mathsf{com}(f(X))$ and $f(a) = b$. PCSs form an essential part of most modern Succinct Non-interactive ARguments of Knowledge (SNARKs). They allow a protocol designer to focus on designing a so-called polynomial Interactive Oracle Proof which can then be compiled, via a PCS, to a SNARK (see [BFS19, GWC19, CHM+19] for descriptions of such compilers). In fact,

many of the most important properties of a SNARK, like proof size, verifier complexity, and cryptographic assumptions, follow primarily from the PCS. The earliest polynomial commitment schemes [KZG10] supported univariate polynomials and were used to construct SNARKs like Plonk [GWC19] and Marlin [CHM$^+$19] with $O(n \log n)$ prover complexity and constant proof size. A different class of SNARKs [Set19, CBBZ22] arising from the sumcheck protocol [LFKN92] have linear prover time, but require *multilinear* Polynomial Commitment Schemes (ml-PCS's).

1.1 Our results

Existing transformations from a univariate PCS to ml-PCS are either linear time but require a logarithmic opening proof size, like gemini [BCHO22] and zeromorph [KT23], or have constant size opening proofs but incur an additional $O(n \log n)$ prover cost to perform univariate polynomial multiplication via FFT's. We propose a new protocol that goes beyond this tradeoff: 𝔐𝔢𝔯𝔠𝔲𝔯𝔶 has constant proof size and only $O(n)$ prover operations (in addition to the $O(\lambda n/(\log(\lambda n))$ operations for muti-scalar multiplications arising in KZG commitments). It is also concretely more efficient than existing schemes with similar verifier complexity[1]in terms of the required scalar multiplications, as can be seen in table 1.

Table 1: Comparison of pairing-based ml-PCS. $\mathbb{G}$ denotes a scalar multiplication. All verifiers below additionally require two pairings. Proof size is measured in elements of $\mathbb{F}$, and uses the fact that a $\mathbb{G}$-element is encoded by two $\mathbb{F}$-elements.

Scheme	Proof size	Prover Work	Verifier Work
univariate-based e.g.[PH23]	$O(1)$	$O(n \log n)$ $\mathbb{F}$, $O(n)$ $\mathbb{G}$	$O(\log n)$ $\mathbb{F}$, $O(1)$ $\mathbb{G}$
gemini [BCHO22]	$O(\log n)$	$O(n)$ $\mathbb{F}$, $3n$ $\mathbb{G}$	$O(\log n)$ $\mathbb{F}$, $O(\log n)$ $\mathbb{G}$
zeromorph [KT23]	$O(\log n)$	$O(n)$ $\mathbb{F}$, $2.5n$ $\mathbb{G}$	$O(\log n)$ $\mathbb{F}$, $O(\log n)$ $\mathbb{G}$
𝔐𝔢𝔯𝔠𝔲𝔯𝔶 (this work)	$O(1)$	$O(n)$ $\mathbb{F}$, $2n + O(\sqrt{n})$ $\mathbb{G}$	$O(\log n)$ $\mathbb{F}$, $O(1)$ $\mathbb{G}$

2 Overview of technique

In this overview, we use some of the notation defined in Sections 3.1 and 3.2.

Our technique is best thought of as an improvement of the gemini ml-PCS [BCHO22]. Let's start by recalling how gemini works. gemini commits to a multilinear function as a *univariate* KZG commitment[KZG10]. Specifically, fix a vector $f \in \mathbb{F}^n$ describing the function's values on the boolean cube $\mathbf{B}_s$ where $s = \log n$. That is, we think of f as

---

[1]The ml-PCS from [KZHB25] requires only $O(\sqrt{n})$ scalar multiplications. However, it requires $O(\log n)$ proof length and $O(\log n)$ verifier pairings whereas all schemes in table 1 require only two.

representing the multilinear

$$M(X_0, \ldots, X_{s-1}) = \sum_{i<n} \mathbf{eq}(i, X_0, \ldots, X_{s-1}) f_i.$$

(Here, as explained in Section 3.2, we interpret i as its binary decomposition $(i_0, \ldots, i_{s-1})$ when used as input to $\mathbf{eq}$.) Let $\mathsf{srs} = \left\{ \left[x^i \right] \right\}_{i<n}$ be a KZG structured reference string. gemini outputs $\mathsf{cm} = [f(x)] = \sum_{i<n} f_i \left[x^i \right]$ as a commitment to M.

Now suppose prover $\mathbf{P}$ wants to convince verifier $\mathbf{V}$ that $M(z) = v$, for some $z = (z_0, \ldots, z_{s-1}) \in \mathbb{F}^s$. In gemini, $\mathbf{P}$ sends commitments $\mathsf{cm}_1, \ldots, \mathsf{cm}_s$ to the s incremental restrictions leading to evaluation at z. Namely, to $M_1 = M(z_0, X_1, \ldots, X_{s-1})$, $M_2 = M(z_0, z_1, X_2, \ldots, X_{s-1}), \ldots, M_s = M(z_0, \ldots, z_{s-1})$. Assuming $\mathbf{P}$ sent commitments to the correct functions, all that is needed is to check that cm_s is the commitment to the constant v. Of course, the interesting part is proving the commitments *are* to the correct functions!

For this purpose, gemini exploits a connection between M and its corresponding univariate $f(X)$: Write $f(X) = f_0(X^2) + X f_1(X^2)$, for $f_0(X), f_1(X)$ of degree $< n/2$. Let $f_{z_0}(X)$ be the univariate corresponding to M_1 defined above. Then, we have

$$f_{z_0}(X) = (1 - z_0) f_0(X) + z_0 f_1(X).$$

Additionally, we can evaluate f_0 and f_1 via f using the equations

$$f_0(X^2) = \frac{f(X) + f(-X)}{2}, f_1(X^2) = \frac{f(X) - f(-X)}{2X}$$

Thus, we can perform consistency checks between each pair $\mathsf{cm}_{i-1}, \mathsf{cm}_i$, via univariate KZG openings at a random challenge, inductively showing cm_i is indeed the commitment to the next desired restriction. Of course, we get $O(s) = O(\log n)$ proof length due to this sequence of restriction commitments.

Here is a first idea on how to reduce proof length. Protocols based on univariate polynomials allow us to do multilinear evaluation in $O(n \log n)$ prover time with constant proof size (e.g. Section 5 of [PH23]). Choose a parameter t and set $b = 2^t$. We can run *only the first t rounds* of gemini, reaching a restricted multilinear on $n - t$ variables. If $n' = n/b \leq n/\log n$, we can afford to run a univariate protocol with $O(n' \log n') = O(n)$ prover time to evaluate $M_t(z_t, \ldots, z_{s-1})$. This still doesn't take us to overall constant proof size - as we need to use a super-constant t to reach such n'. (For us $t = \log n/2$ will be optimal, although $t \geq \log \log n$ suffices here.)

This raises the question - can we "skip" the intermediate gemini rounds and send *only* the commitment cm_t, and directly prove it is consistent with the original cm? Extrapolating the gemini strategy in the natural way, we get the answer - yes, but not with constant proof size: We can decompose f into b polynomials of degree $< n/b$: $f(X) = \sum_{0 \leq i < b} X^i f_i(X^b)$. As in the $b = 2$ case, one can show the univariate $f'(X)$ corresponding to M_t is a linear combination of the $\{f_i(X)\}$. Moreover, evaluating the f_i using f (for the consistency check) can be done. However, it requires b evaluations of

3

f. Specifically, $f_i(r^b)$ is a linear combination of $\left\{ f(r), f(r\omega), \ldots, f(r\omega^{b-1}) \right\}$ where ω is a primitive b'th root of unity.

Our central innovation is a different way to prove cm_t is correct *with* constant proof size. Let's switch notation and denote the opening point as $u = (u_1, u_2)$ where $u_1 \in \mathbb{F}^t, u_2 \in \mathbb{F}^{s-t}$. The (univariate corresponding to the) correct restricted polynomial is

$$h(X) = \sum_{0 \leq i < b} \mathbf{eq}(i, u_1) f_i(X).$$

Let $g(X) := f(X) \bmod X^b - \alpha$. Calculation shows

$$g(X) = \sum_{0 \leq i < b} X^i f_i(\alpha).$$

The multilinear $\hat{g}$ corresponding to $g(X)$ is

$$\hat{g}(X_0, \ldots X_{t-1}) = \sum_{0 \leq i < b} \mathbf{eq}(i, X_0, \ldots, X_{t-1}) f_i(\alpha).$$

In particular, we have

$$\hat{g}(u_1) = \sum_{0 \leq i < b} \mathbf{eq}(i, u_1) f_i(\alpha) = h(\alpha).$$

In words, the evaluation of g at u_1 as a *multilinear* corresponds to the evaluation of h at α as a univariate! We can use standard univariate KZG to open cm_t at α. And, crucially, we can afford to evaluate $\hat{g}(u_1)$ using the aforementioned univariate protocols as it is of size b rather than n. In summary, we can show a committed polynomial corresponds to the correct restriction. And now, again, we can afford to open h as a multilinear at u_2 using univariate protocols as it has size n/b rather than n.

Comparison to sumcheck   It is instructive to see what happens if we try to get a similar result via a modification of the sumcheck protocol [LFKN92]. Note first that a multilinear evaluation can indeed be written as a sum over the function's values on $\mathbf{B}_s$ multiplied by the $\mathbf{eq}$ function:

$$M(u) = \sum_{b \in \mathbf{B}_s} \mathbf{eq}(b, u) M(b).$$

The classic sumcheck protocol, like gemini, works by $\log n$ reductions of the domain size by a factor of two; each round fixing one more variable of the summed function. In the above spirit, we could look at a modified sumcheck protocol, where the first variable ranges over a domain of super-constant size b. The first round univariate P_1 would thus have degree roughly $2b$. To maintain constant proof length, we could send a commitment to P_1 rather than its coefficients (as usually done in sumcheck). However, *computing* P_1 would require superlinear time $O(n \log b)$ - as we need to perform a b-size FFT for n/b values of the second variable appearing in the sum.

3 Preliminaries

3.1 Terminology and conventions

Fields and Groups We assume our field $\mathbb{F}$ is of prime order. We denote by $\mathbb{F}_{<d}[X]$ the set of univariate polynomials over $\mathbb{F}$ of degree smaller than d. We assume all algorithms described receive as an implicit parameter the security parameter λ.

Whenever we use the term *efficient*, we mean an algorithm running in time $\mathsf{poly}(\lambda)$. Furthermore, we assume an *object generator* $\mathcal{O}$ that is run with input λ before all protocols, and returns all fields and groups used. Specifically, in our protocol $\mathcal{O}(\lambda) = (\mathbb{F}, \mathbb{G}, \mathbb{G}_2, \mathbb{G}_t, e, \mathcal{g}, \mathcal{g}_2, \mathcal{g}_t)$ where

- $\mathbb{F}$ is a prime field of super-polynomial size $r = \lambda^{\omega(1)}$.

- $\mathbb{G}, \mathbb{G}_2, \mathbb{G}_t$ are all groups of size r, and e is an efficiently computable non-degenerate pairing $e : \mathbb{G} \times \mathbb{G}_2 \to \mathbb{G}_t$.

- $\mathcal{g}, \mathcal{g}_2$ are uniformly chosen generators such that $e(\mathcal{g}, \mathcal{g}_2) = \mathcal{g}_t$.

We usually let the λ parameter be implicit, i.e. write $\mathbb{F}$ instead of $\mathbb{F}(\lambda)$. We write $\mathbb{G}$ and $\mathbb{G}_2$ additively. We use the notations $[x] := x \cdot \mathcal{g}$ and $[x]_2 := x \cdot \mathcal{g}_2$.

Vectors and polynomials We work with integer parameter n that we'll assume throughout the paper is of the form $n = 2^{2t}$ for integer $t > 0$. We'll denote its square root by $b := 2^t = \sqrt{n}$. We index vectors starting from zero. For example, for $g \in \mathbb{F}^b$ we have $g = (g_0, \ldots, g_{b-1})$. We associate vectors with univariate polynomials in the following natural way: Given $g \in \mathbb{F}^b$ we denote $g(X) := \sum_{0 \le i < b} g_i X^i$.

We make the convention that integer ranges in sums begin at zero if not specified otherwise. Thus, we write $g(X) = \sum_{i<b} g_i X^i$.

We assume vectors of size n are indexed by two indices ranging over $\{0, \ldots, b-1\}$. Thus, for $f \in \mathbb{F}^n$, we have $f = (f_{0,0}, \ldots, f_{0,b-1}, \ldots, f_{b-1,0}, \ldots, f_{b-1,b-1})$. For $0 \le i < b$, we denote by f_i the vector $(f_{0,i}, \ldots, f_{b-1,i})$.

In particular, for $f \in \mathbb{F}^n$ we have under these notations that

$$f(X) := \sum_{i<b} X^i f_i(X^b) = \sum_{i<b} \sum_{j<b} f_{i,j} X^{i+j\cdot b}$$

For integet $m > 0$, we denote by $\mathbf{B}_m$ the binary cube $\{0,1\}^m \subset \mathbb{F}^m$ of dimension m.

3.2 Multilinear polynomials

Let $n = 2^{2t}$, and $s = 2t$. We define the well-known **eq** multilinear polynomial in $2s$ variables.

$$\mathbf{eq}(x, y) := \prod_{i=0}^{s-1} (x_i y_i + (1 - x_i)(1 - y_i))$$

We have for $x, y \in \mathbf{B}_s$, $\mathbf{eq}(x, y) = 1$ when $x = y$ and $\mathbf{eq}(x, y) = 0$ otherwise.

5

We use the convention that an integer $0 \leq i < n$ can be used as an input to **eq** by interpreting i as its binary representation. Namely, for $0 \leq i < n$, $u \in \mathbb{F}^s$, $\mathbf{eq}(i, u) := \mathbf{eq}(i_0, \ldots, i_{s-1}, u)$ where $i = \sum_{j<s} i_j 2^j$.

For $f \in \mathbb{F}^n$, we define $\hat{f}$ to be the multilinear polynomial obtaining f's values on the boolean cube. Namely,

$$\hat{f}(X_0, \ldots, X_{s-1}) := \sum_{i<n} \mathbf{eq}(i, X_0, \ldots, X_{s-1}) \cdot f_i.$$

Decomposing eq We'll overload **eq** to also denote the analogous equality function for $x, y \in \mathbf{B}_t$. With this overloading, given $w_1, w_2, u_1, u_2 \in \mathbf{B}_t$ we have the convenient decomposition

$$\mathbf{eq}((w_1, w_2), (u_1, u_2)) = \mathbf{eq}(w_1, u_1)\mathbf{eq}(w_2, u_2).$$

3.3 The algebraic group model

We introduce some terminology from [GWC19] to capture analysis in the Algebraic Group Model of Fuchsbauer, Kiltz and Loss[FKL18]. In this subsection, we use the notation $\mathbb{G}_1 = \mathbb{G}$. In our protocols, by an *algebraic adversary* $\mathcal{A}$ in an SRS-based protocol we mean a $\mathsf{poly}(\lambda)$-time algorithm which satisfies the following.

- For $i \in \{1, 2\}$, whenever $\mathcal{A}$ outputs an element $A \in \mathbb{G}_i$, it also outputs a vector v over $\mathbb{F}$ such that $A = < v, \mathsf{srs}_i >$.

First we say our srs *has degree* Q if all elements of srs_i are of the form $[f(x)]_i$ for $f \in \mathbb{F}_{<Q+1}[X]$ and uniform $x \in \mathbb{F}$. In the following discussion let us assume we are executing a protocol with a degree Q SRS, and denote by $f_{i,j}$ the corresponding polynomial for the j'th element of srs_i.

Denote by a, b the vectors of $\mathbb{F}$-elements whose encodings in $\mathbb{G}_1, \mathbb{G}_2$ an algebraic adversary $\mathcal{A}$ outputs during a protocol execution; e.g., the j'th $\mathbb{G}_1$ element output by $\mathcal{A}$ is $[a_j]$.

By a "real pairing check" we mean a check of the form

$$(a \cdot T_1) \cdot (T_2 \cdot b) = 0$$

for some matrices T_1, T_2 over $\mathbb{F}$. Note that such a check can indeed be done efficiently given the encoded elements and the pairing function $e : \mathbb{G}_1 \times \mathbb{G}_2 \to \mathbb{G}_t$.

Given such a "real pairing check", and the adversary $\mathcal{A}$ and protocol execution during which the elements were output, define the corresponding "ideal check" as follows. Since $\mathcal{A}$ is algebraic when he outputs $[a_j]_i$ he also outputs a vector v such that, from linearity, $a_j = \sum v_\ell f_{i,\ell}(x) = R_{i,j}(x)$ for $R_{i,j}(X) := \sum v_\ell f_{i,\ell}(X)$. Denote, for $i \in \{1, 2\}$ the vector of polynomials $R_i = (R_{i,j})_j$. The corresponding ideal check, checks as a polynomial identity whether

$$(R_1 \cdot T_1) \cdot (T_2 \cdot R_2) \equiv 0$$

6

The following lemma is inspired by [FKL18]'s analysis of [Gro16], and tells us that for soundness analysis against algebraic adversaries it suffices to look at ideal checks. Before stating the lemma we define the Q-DLOG assumption similarly to [FKL18].

Definition 3.1. *Fix integer Q. The Q-DLOG assumption for $(\mathbb{G}_1, \mathbb{G}_2)$ states that given*

$$[1], [x], \ldots, \left[x^Q\right], [1]_2, [x]_2, \ldots, \left[x^Q\right]_2$$

for uniformly chosen $x \in \mathbb{F}$, the probability of an efficient $\mathcal{A}$ outputting x is $\mathsf{negl}(\lambda)$.

Lemma 3.2. *Assume the Q-DLOG for $(\mathbb{G}_1, \mathbb{G}_2)$. Given an algebraic adversary $\mathcal{A}$ participating in a protocol with a degree Q SRS, the probability of any real pairing check passing is larger by at most an additive $\mathsf{negl}(\lambda)$ factor than the probability the corresponding ideal check holds.*

See [GWC19] for the proof.

3.4 Polynomial commitment schemes for multilinear polynomials

We give a formal definition of an ml-PCS secure in the algebraic group model.

Definition 3.3. *Let $n = 2^s$. A multilinear polynomial commitment scheme (ml-PCS) consists of*

- $\mathsf{gen}(n)$ *- a randomized algorithm that outputs an SRS* srs.

- $\mathsf{com}(f, \mathsf{srs})$ *- that given a polynomial $f \in \mathbb{F}^n$ returns a commitment cm to f.*

- *A public coin protocol $\mathsf{open}(\mathsf{cm}, n, u, v)$ between parties $\mathbf{P}$ and $\mathbf{V}$. $\mathbf{P}$ is given $f \in \mathbb{F}^n$. $\mathbf{P}$ and $\mathbf{V}$ are both given integer n, cm- the purported commitment to f, $u \in \mathbb{F}^s$ and $v \in \mathbb{F}$ - the purported value $\hat{f}(u)$.*

such that

- ***Completeness:*** *Suppose that $\mathsf{cm} = \mathsf{com}(f, \mathsf{srs})$. Then if open is run correctly with values $n, \mathsf{cm}, u, v = \hat{f}(u)$, $\mathbf{V}$ outputs $accept$ with probability one.*

- ***Knowledge soundness in the algebraic group model:*** *There exists an efficient $\mathcal{E}$ such that for any efficient algebraic adversary $\mathcal{A}$ the probability of $\mathcal{A}$ winning the following game is $\mathsf{negl}(\lambda)$ over the randomness of $\mathcal{A}$ and gen.*

    1. *Given srs, $\mathcal{A}$ outputs n, cm.*
    2. *$\mathcal{E}$, given access to the messages of $\mathcal{A}$ during the previous step, outputs $f \in \mathbb{F}^n$.*
    3. *$\mathcal{A}$ outputs $u \in \mathbb{F}^s$ and $v \in \mathbb{F}$.*
    4. *$\mathcal{A}$ takes the part of $\mathbf{P}$ in the protocol open with inputs n, cm, u, v.*
    5. *$\mathcal{A}$ wins if*
        - *$\mathbf{V}$ outputs $accept$ at the end of the protocol.*
        - *$\hat{f}(u) \neq v$.*

4 Components

In this section we go over known components (with some new optimizations), that will be used in our main protocol in Section 6. The treatment will be semi-formal, and assume basic familiarity with the KZG polynomial commitment scheme [KZG10]. The formal treatment will be part of the description and knowledge soundness proof of the main protocol in Section 6.

4.1 Inner products in $O(b \log b)$ time.

Fix polynomials $g_1(X) = \sum_{i=0}^{d_1} a_i X^i, g_2 = \sum_{i=0}^{d_2} b_i X^i$ in $\mathbb{F}[X]$. We define $< g_1, g_2 >$ to be $\sum_{i=0}^{d} a_i b_i$ where $d := \min\{d_1, d_2\}$. We present a convenient way to verify inner products $< g_1, g_2 >$ similar to [BCC+16, MBKM19]. The basic observation is that $< g_1, g_2 >$ is the constant coefficient of the rational function $R(X) := g_1(X)g_2(1/X)$. Thus, $< g_1, g_2 > = v$ is equivalent to the existence of polynomials $S_1(X), S_2(X)$ such that

$$g_1(X)g_2(1/X) = 1/X \cdot S_1(1/X) + v + X \cdot S_2(X).$$

We can thus sends commitments to S_1, S_2 as proof of the correctness of v. As an optimization, we observe that we can "symmetrize" R and look instead at the rational function

$$R'(X) := g_1(X)g_2(1/X) + g_1(1/X)g_2(X).$$

The advantage of R' is that the negative and positive coefficients are equal. Thus, $< g_1, g_2 > = v$ is equivalent to the existence of $S(X) \in \mathbb{F}[X]$ such that

$$g_1(X)g_2(1/X) + g_1(1/X)g_2(X) = 2v + X \cdot S(X) + (1/X)S(1/X).$$

Claim 4.1. *Suppose $g_1(X), g_2(X) \in \mathbb{F}_{<b}[X]$. Let $S(X)$ be as defined above. Then S can be computed in $O(b \log b)$ $\mathbb{F}$-operations.*

Proof. When $g_1(X), g_2(X) \in \mathbb{F}_{<b}[X]$ we multiply the equation above by X^{b-1} to get

$$X^{b-1}(g_1(X)g_2(1/X) + g_1(1/X)g_2(X)) = X^{b-1}(2v + X \cdot S(X) + (1/X)S(1/X)).$$

We can use an $O(b \log b)$ time FFT to evaluate the LHS on $2b$ points. We then do an inverse FFT to get the coefficients $c_0, \ldots, c_{2b-2}$ of the LHS. Now, we can output $S = (c_b, \ldots, c_{2b-2})$. $\qquad\square$

Batching inner product checks Suppose we now have two inner product claims $< g_1, g_2 > = v_1$ and $< h_1, h_2 > = v_2$. The following claim gives us a way to randomly batch them so that one polynomial $S(X)$ suffices to prove both.

Claim 4.2. *Fix polynomials $g_1, g_2, h_1, h_2 \in \mathbb{F}[X]$ and $v_1, v_2 \in \mathbb{F}$. Suppose $< g_1, g_2 > \neq v_1$ or $< h_1, h_2 > \neq v_2$. Then, e.w.p $1/|\mathbb{F}|$ over $\gamma \in \mathbb{F}$ there does not exist $S(X) \in \mathbb{F}[X]$ such that*

$$g_1(X)g_2(1/X) + g_1(1/X)g_2(X) + \gamma(h_1(X)h_2(1/X) + h_1(1/X)h_2(X))$$

$$= 2(v_1 + \gamma v_2) + X \cdot S(X) + (1/X)S(1/X)$$

Proof. Denote $(v_1', v_2') = (< g_1, g_2 >, < h_1, h_2 >))$. The constant coefficient of the LHS is $v_1' + \gamma v_2'$. To satisfy the equation in the claim for some S, we need

$$v_1' + \gamma v_2' = v_1 + \gamma v_2,$$

which can hold for at most one γ when $(v_1, v_2) \neq (v_1', v_2')$. $\square$

4.2 Multilinear evaluations as inner products of univariate polynomials.

For $u \in \mathbb{F}^t$ define the polynomial $P_u(X) := \sum_{i<b} \mathbf{eq}(i, u)X^i$. Note that for $g(X) \in \mathbb{F}_{<b}[X]$, we have

$$< P_u, g >= \sum_{i<b} \mathbf{eq}(i, u)g_i = \hat{g}(u).$$

As leveraged in [BGH19], we have the product formula

$$P_u(X) = \prod_{i=0}^{t-1} \left(u_i X^{2^i} + 1 - u_i \right),$$

implying $P_u(X)$ can be evaluated in $O(t)$ $\mathbb{F}$-operations. Hence, a verifier $\mathbf{V}$ operating in $O(\log n)$ time can evaluate $P_u(X)$ itself.

Using the polynomials P_u, mutlilinear evaluations can be proven in a batched manner based on Claim 4.2: Suppose we want to show given committed univariates $g(X), h(X) \in \mathbb{F}[X]$, $u_1, u_2 \in \mathbb{F}^t, v_1, v_2 \in \mathbb{F}$ that $\hat{g}(u_1) = v_1$ and $\hat{h}(u_2) = v_2$.

1. $\mathbf{V}$ sends random $\gamma \in \mathbb{F}$.

2. $\mathbf{P}$ sends commitment to S such that

$$g(X)P_{u_1}(1/X) + g(1/X)P_{u_1}(X) + \gamma(h(X)P_{u_2}(1/X) + h(1/X)P_{u_2}(X))$$

$$= 2(v_1 + \gamma v_2) + X \cdot S(X) + (1/X)S(1/X).$$

3. $\mathbf{V}$ chooses a random $\mathfrak{z} \in \mathbb{F}$.

4. $\mathbf{P}$ sends and proves correctness of the values of g, h and S on $\mathfrak{z}, 1/\mathfrak{z}$.

5. $\mathbf{V}$ evaluates P_{u_1}, P_{u_2} at $\mathfrak{z}, 1/\mathfrak{z}$.

6. $\mathbf{V}$ checks the equation in step 2 holds at $\mathfrak{z}$.

9

4.3 Degree checks

The idea presented here is from [Tha23]. Suppose $\mathbf{P}$ wants to prove to $\mathbf{V}$ that cm is a commitment to a polynomial $g(X) \in \mathbb{F}_{<b}[X]$. Let $D(X) := X^{b-1}g(1/X)$. The idea is that $D(X)$ is a polynomial if and only if $g(X)$ has degree $< b$. Thus, assuming our structured reference string doesn't contain negative powers, $\mathbf{P}$ can commit to D if and only if $g(X) \in \mathbb{F}_{<b}[X]$.

This motivates the following protocol.

1. $\mathbf{P}$ sends a commitment d to $D(X)$.

2. $\mathbf{V}$ chooses random $\mathfrak{z} \in \mathbb{F}$.

3. $\mathbf{P}$ sends $D_{\mathfrak{z}} := D(\mathfrak{z}), \bar{g}_{\mathfrak{z}} := g(1/\mathfrak{z})$, and uses KZG to prove their correctness.

4. $\mathbf{V}$ can now check D's correctness on $\mathfrak{z}$, using the equation

$$D_{\mathfrak{z}} \overset{?}{=} \mathfrak{z}^{b-1}\bar{g}_{\mathfrak{z}}.$$

5 Univariate division

Our protocol crucially relies on the following simple claim about division by a polynomial of the form $X^b - \alpha$.

Claim 5.1. *Fix integers $b > 0$ and let $n = b^2$. Fix $\alpha \in \mathbb{F}$, and $f(X) \in \mathbb{F}_{<n}[X]$. Let $f_0(X), \ldots, f_{b-1}(X) \in \mathbb{F}_{<b}[X]$ be such that $f(X) = \sum_{i<b} X^i f_i(X^b)$. Let $g(X) \in \mathbb{F}_{<b}[X], q(X) \in \mathbb{F}[X]$ be such that*

$$f(X) = (X^b - \alpha) \cdot q(X) + g(X).$$

Then,

1. $g(X) = \sum_{i<b} X^i f_i(\alpha)$.

2. The coefficients of $q(X)$ can be computed in $O(n)$ $\mathbb{F}$-operations.

Proof. To see the first item, note that reduction mod $X^b - \alpha$ corresponds to substituting α into X^b inside each $f_i(X^b)$ in the expression $\sum_{i<b} X^i f_i(X^b)$. We proceed to the computation of $q(X)$. We compute for each $0 \le i < b$, the coefficients of the quotient $q_i(X) \in \mathbb{F}[X]$ such that
$$f_i(X) = q_i(X)(X - \alpha) + f_i(\alpha).$$

Using Horner's method for division by the linear polynomial $X - \alpha$ this requires only n multiplications and additions in $\mathbb{F}$. Now, we have that

$$f(X) = \sum_{i<b} X^i f_i(X^b) = \sum_{i<b} X^i \left(q_i(X^b)(X^b - \alpha) + f_i(\alpha) \right) = q(X)(X^b - \alpha) + g(X),$$

for $q(X) := \sum_{i<b} X^i q_i(X^b)$. Thus, the coefficients of $q(X)$ are simply the interleaving of the coefficients of the $\{q_i(X)\}$. $\qquad\square$

6 Main Construction

𝔐𝔈𝔯𝔠𝔲𝔯𝔶 is the tuple (gen, com, open) described next.

gen(n): Choose random $x \in \mathbb{F}$ and outputs $\left\{[1], [x], \ldots, [x^{n-1}], [1]_2, [x]_2\right\}$

com(n, f, srs): Output $\sum_{i<b} \sum_{j<b} f_{i,j} \cdot [x^{i+j\cdot b}]$.

open($n, \mathsf{cm}, u, v; f$):

1. Committing to partial sums:

   (a) Let $u = (u_1, u_2)$ for $u_1, u_2 \in \mathbb{F}^t$. **P** computes the polynomial $h(X) := \sum_{i<b} \mathbf{eq}(i, u_1) f_i(X)$. Note that the coefficient of X^j in $h(X)$ is $\sum_{i<b} \mathbf{eq}(i, u_1) f_{i,j}$ - hence we think of it as a commitment to partial sums.

   (b) **P** computes and sends $\mathsf{h} := [h(x)]$.

2. Committing to "folded" polynomial g:

   (a) **V** sends random $\alpha \in \mathbb{F}$.

   (b) **P** computes polynomials $g(X) \in \mathbb{F}_{<b}[X]$ and $q(X) \in \mathbb{F}[X]$ such that

   $$f(X) = (X^b - \alpha) \cdot q(X) + g(X).$$

   (c) **P** computes and sends $\mathsf{q} := [q(x)]$ and $\mathsf{g} := [g(x)]$.

3. Sending proofs of correctness for h and the degree of g:

   (a) **V** sends a random batching challenge $\gamma \in \mathbb{F}$.

   (b) **P** computes and sends $\mathsf{s} = [S(x)]$ where $S(X) \in \mathbb{F}[X]$ is such that

   $$g(X)P_{u_1}(1/X) + g(1/X)P_{u_1}(X) + \gamma \cdot (h(X)P_{u_2}(1/X) + h(1/X)P_{u_2}(X))$$

   $$= 2(h(\alpha) + \gamma \cdot v) + X \cdot S(X) + (1/X)S(1/X).$$

   (c) **P** computes and sends $\mathsf{d} := [D(x)]$ where

   $$D(X) := X^{b-1}g(1/X).$$

4. KZG evaluations:

   (a) **V** sends a random evaluation challenge $\mathfrak{z} \in \mathbb{F}$.

   (b) **P** sends the values $g_{\mathfrak{z}} := g(\mathfrak{z}), \bar{g}_{\mathfrak{z}} := g(1/\mathfrak{z}), h_{\mathfrak{z}} := h(\mathfrak{z}), \bar{h}_{\mathfrak{z}} := h(1/\mathfrak{z}), s_{\mathfrak{z}} := S(\mathfrak{z}), \bar{s}_{\mathfrak{z}} := S(1/\mathfrak{z}).$

(c) **V** computes the expected values for $D(\mathfrak{z})$ and $h(\alpha)$ assuming the equations in steps 3b,3c are satisfied. That is, $D_{\mathfrak{z}} := \mathfrak{z}^{b-1}\bar{g}_{\mathfrak{z}}$, and

$$h_\alpha := \left(g_{\mathfrak{z}}P_{u_1}(1/\mathfrak{z}) + \bar{g}_{\mathfrak{z}}P_{u_1}(\mathfrak{z}) + \gamma(h_{\mathfrak{z}}P_{u_2}(1/\mathfrak{z}) + \bar{h}_{\mathfrak{z}}P_{u_2}(\mathfrak{z}) - 2v) - \mathfrak{z}s_{\mathfrak{z}} - (1/\mathfrak{z})\bar{s}_{\mathfrak{z}}\right)/2.$$

(d) **P** computes and sends the KZG opening proof $\pi_{\mathfrak{z}}$ to check the equation of step 2b at $\mathfrak{z}$. That is $\pi_{\mathfrak{z}} := [H(x)]$ for

$$H(X) := \frac{f(X) - (\mathfrak{z}^b - \alpha)q(X) - g_{\mathfrak{z}}}{X - \mathfrak{z}}.$$

(e) **P** computes and sends a batched KZG opening proof π' for the values sent in step 4b and computed by **V** in step 4c, as described in Section 4 of [BDFG20].

(f) **V** checks the proof $\pi_{\mathfrak{z}}$ via pairings as in [KZG10]:

$$e(\mathsf{cm} - \left[\mathfrak{z}^b - \alpha\right] \cdot \mathsf{q} - g_{\mathfrak{z}}, [1]_2) = e(\pi_{\mathfrak{z}}, [x - \mathfrak{z}]_2).$$

(g) **V** checks the opening proof π' as described in [BDFG20].

(h) If one of the checks in steps 4f,4g fails, **V** outputs *reject*. Otherwise **V** outputs *accept*.

Runtime of **P**: Computing $q(X)$ in step 2b requires $O(n)$ operations by Claim 5.1. Computing q and $\pi_{\mathfrak{z}}$ requires two MSMs of size n. All other steps are on polynomials of size $O(b) = O(\sqrt{n})$. Thus, other commitments clearly require $O(\sqrt{n})$ scalar multiplications. It is easy to see other steps require $o(n)$ $\mathbb{F}$-operations. The least trivial of these is perhaps the computation of $S(X)$ shown to require $O(b \log b) = o(n)$ operations in Claim 4.1.

Proving knowledge soundness: Let $\mathcal{A}$ be an efficient algebraic adversary participating in the Knowledge Soundness game from Definition 3.3. We show its probability of winning the game is $\mathsf{negl}(\lambda)$. We define the extractor $\mathcal{E}$ to simply output the vector $f \in \mathbb{F}^n$ $\mathcal{A}$ outputs (as it's algebraic) with $\mathsf{com}(f) = \mathsf{cm}$ together with cm.

As $\mathcal{A}$ is algebraic, when sending the commitments $\mathsf{h,q,g,s,d},\pi_{\mathfrak{z}},\pi'$ during protocol execution it also sends polynomials $h(X), q(X), g(X), S(X), D(X), H(X), Q(X) \in \mathbb{F}_{<n}[X]$ such that the former are their corresponding commitments. Let E be the event that **V** outputs *accept*. Let A be the event that $\mathcal{A}$ wins the knowledge soundness game. Note that by definition $A \subset E$, and our goal is to show $\mathsf{prob}(A) = \mathsf{negl}(\lambda)$. We will define a constant number of events such that their union contains A and each has probability $\mathsf{negl}(\lambda)$. This implies the knowledge soundness of the protocol.

E implies all pairing checks have passed. Let $E_0 \subset E$ be the event that one of the corresponding ideal pairing checks as defined in Section 3.3 didn't pass. According to Lemma 3.2, $\mathsf{prob}(E_0) = \mathsf{negl}(\lambda)$.

Given that E_0 didn't occur, we have from the knowledge soundness proof of batched KZG in Section 3 of [BDFG20] that the evaluations sent by **P** and computed by **V** are all correct. That is,

1. $g_{\mathfrak{z}} = g(\mathfrak{z}), \bar{g}_{\mathfrak{z}} = g(1/\mathfrak{z}), h_{\mathfrak{z}} = h(\mathfrak{z}), \bar{h}_{\mathfrak{z}} = h(1/\mathfrak{z}), s_{\mathfrak{z}} = S(\mathfrak{z}), \bar{s}_{\mathfrak{z}} = S(1/\mathfrak{z})$,

2. $D(\mathfrak{z}) = \mathfrak{z}^{b-1}g(1/\mathfrak{z})$,

3.

$2(h(\alpha)+\gamma v) = g(\mathfrak{z})P_{u_1}(1/\mathfrak{z})+g(1/\mathfrak{z})P_{u_1}(\mathfrak{z})+\gamma(h(\mathfrak{z})P_{u_2}(1/\mathfrak{z})+h(1/\mathfrak{z})P_{u_2}(\mathfrak{z}))-\mathfrak{z}S(\mathfrak{z})-(1/\mathfrak{z})S(1/\mathfrak{z})$.

4. $g(\mathfrak{z}) = f(\mathfrak{z}) - (\mathfrak{z}^b - \alpha)q(\mathfrak{z}) - g(\mathfrak{z})$.

Note that items 2-4 can be viewed as rational equations evaluated at $\mathfrak{z}$. Let E_1 be the event that E_0 didn't occur, and one of the equations in steps 2-4 doesn't hold as a rational identity. Multiplying denominators, we have from the Schwarz-Zippel Lemma that E_1 occurs with probability at most $2n/|\mathbb{F}| = \mathsf{negl}(\lambda)$ over $\mathfrak{z} \in \mathbb{F}$. Assuming E_0 and E_1 didn't occur we have that

1. $D(X) = X^{b-1}g(1/X)$,

2.

$2(h(\alpha)+\gamma v) = g(X)P_{u_1}(1/X)+g(1/X)P_{u_1}(X)+\gamma(h(X)P_{u_2}(1/X)+h(1/X)P_{u_2}(X))$

$$-XS(X) - (1/X)S(1/X).$$

3. $g(X) = f(X) - (X^b - \alpha)q(X)$.

Let E_2 be the event that E_0 and E_1 don't occur but $\hat{g}(u_1) \neq h(\alpha)$ or $\hat{h}(u_2) \neq v$. According to Claim 4.2, given the equation in item 2, E_2 occurs with probability at most $1/|\mathbb{F}|$ over γ.

Let E_3 be the event that $E_0 \cup E_1 \cup E_2$ don't occur and

1. $h(X) \neq \sum_{i<b} \mathbf{eq}(i, u_1)f_i(X)$,

2. $h(\alpha) = \sum_{i<b} \mathbf{eq}(i, u_1)f_i(\alpha)$.

Obviously, E_3 has probability $\mathsf{negl}(\lambda)$.

Assume $E_0 \cup E_1 \cup E_2 \cup E_3$ doesn't occur. We show that $\hat{f}(u) = v$, and thus we are outside the event A. In other words, $A \subset E_0 \cup E_1 \cup E_2 \cup E_3$.

Since $D(X) = X^{b-1}g(1/X)$ we know that $\deg(g) < b$. Since $g(X) = f(X) - (X^b - \alpha)q(X)$, from Claim 5.1 we know that $g(X) = \sum_{i<b} f_i(\alpha)X^i$. Hence, we know that

$$\hat{g}(u_1) = \sum_{i<b} \mathbf{eq}(i, u_1)f_i(\alpha) = h(\alpha).$$

Using $\neg E_3$ we know that $h(X) = \sum_{i<b} \mathbf{eq}(i, u_1)f_i(X)$. Hence, writing $h(X) = \sum_{j<b} h_j X^j$ we have $h_j = \sum_{i<b} \mathbf{eq}(i, u_1)f_{i,j}$. Thus, we have

$$\hat{f}(u) = \sum_{i<b}\sum_{j<b} \mathbf{eq}(i, u_1)\mathbf{eq}(j, u_2)f_{i,j} = \sum_{j<b} \mathbf{eq}(j, u_2)\sum_{i<b} \mathbf{eq}(i, u_1)f_{i,j}$$

$$= \sum_{j<b} \mathbf{eq}(j, u_2)h_j = \hat{h}(u_2) = v.$$

13

Acknowledgements

References

[BCC⁺16]   J. Bootle, A. Cerulli, P. Chaidos, J. Groth, and C. Petit. Efficient zero-knowledge arguments for arithmetic circuits in the discrete log setting. pages 327–357, 2016.

[BCHO22]   J. Bootle, A. Chiesa, Y. Hu, and M. Orrù. Gemini: Elastic snarks for diverse environments. *IACR Cryptol. ePrint Arch.*, page 420, 2022.

[BDFG20]   D. Boneh, J. Drake, B. Fisch, and A. Gabizon. Efficient polynomial commitment schemes for multiple points and polynomials. *IACR Cryptol. ePrint Arch.*, page 81, 2020.

[BFS19]   B. Bünz, B. Fisch, and A. Szepieniec. Transparent snarks from DARK compilers. *IACR Cryptol. ePrint Arch.*, page 1229, 2019.

[BGH19]   S. Bowe, J. Grigg, and D. Hopwood. Halo: Recursive proof composition without a trusted setup. *IACR Cryptol. ePrint Arch.*, page 1021, 2019.

[CBBZ22]   B. Chen, B. Bünz, D. Boneh, and Z. Zhang. Hyperplonk: Plonk with linear-time prover and high-degree custom gates. *IACR Cryptol. ePrint Arch.*, page 1355, 2022.

[CHM⁺19]   A. Chiesa, Y. Hu, M. Maller, P. Mishra, N. Vesely, and N. P. Ward. Marlin: Preprocessing zksnarks with universal and updatable SRS. *IACR Cryptology ePrint Archive*, 2019:1047, 2019.

[FKL18]   G. Fuchsbauer, E. Kiltz, and J. Loss. The algebraic group model and its applications. In *Advances in Cryptology - CRYPTO 2018 - 38th Annual International Cryptology Conference, Santa Barbara, CA, USA, August 19-23, 2018, Proceedings, Part II*, pages 33–62, 2018.

[Gro16]   J. Groth. On the size of pairing-based non-interactive arguments. In *Advances in Cryptology - EUROCRYPT 2016 - 35th Annual International Conference on the Theory and Applications of Cryptographic Techniques, Vienna, Austria, May 8-12, 2016, Proceedings, Part II*, pages 305–326, 2016.

[GWC19]   A. Gabizon, Z. J. Williamson, and O. Ciobotaru. PLONK: permutations over lagrange-bases for oecumenical noninteractive arguments of knowledge. *IACR Cryptology ePrint Archive*, 2019:953, 2019.

[KT23]   T. Kohrita and P. Towa. Zeromorph: Zero-knowledge multilinear-evaluation proofs from homomorphic univariate commitments. *IACR Cryptol. ePrint Arch.*, page 917, 2023.

[KZG10]   A. Kate, G. M. Zaverucha, and I. Goldberg. Constant-size commitments to polynomials and their applications. pages 177–194, 2010.

[KZHB25]   G. Kadianakis, A. Zapico, H. Hafezi, and B. Bünz. Kzh-fold: Accountable voting from sublinear accumulation. *IACR Cryptol. ePrint Arch.*, page 144, 2025.

[LFKN92]   C. Lund, L. Fortnow, H. J. Karloff, and N. Nisan. Algebraic methods for interactive proof systems. *J. ACM*, 39(4):859–868, 1992.

[MBKM19] M. Maller, S. Bowe, M. Kohlweiss, and S. Meiklejohn. Sonic: Zero-knowledge snarks from linear-size universal and updateable structured reference strings. *IACR Cryptology ePrint Archive*, 2019:99, 2019.

[PH23]   S. Papini and U. Haböck. Improving logarithmic derivative lookups using GKR. *IACR Cryptol. ePrint Arch.*, page 1284, 2023.

[Set19]   S. T. V. Setty. Spartan: Efficient and general-purpose zksnarks without trusted setup. *IACR Cryptol. ePrint Arch.*, page 550, 2019.

[Tha23]   S. Thakur. A flexible snark via the monomial basis. *IACR Cryptol. ePrint Arch.*, page 788, 2023.